

# GeneralModeling

March 29, 2021

## 1 General Regressions and Modeling

```
[1]: %load_ext autoreload
      %autoreload 2
```

```
[155]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[6]: import os
import sys
HOME_DIR = '/home/aysola/'
os.chdir(HOME_DIR)
sys.path.append('./NVFEL498/')
```

```
[3]: from Notebooks.utils import *
```

```
[9]: data = pd.read_csv('./NVFEL498/alltrips_vel.csv')
```

```
[12]: print(data.shape)
```

(11803, 15)

```
[11]: data.head()
```

```
[11]:
```

|   | Unnamed: 0 | Unnamed: 0.1 | Unnamed: 0.1.1 | TripId | TripId_raw | VehId | \ |
|---|------------|--------------|----------------|--------|------------|-------|---|
| 0 | 0          | 0            | 0              | 0      | 2265       | 123   |   |
| 1 | 1          | 1            | 1              | 1      | 1239       | 135   |   |
| 2 | 2          | 2            | 2              | 4      | 2263       | 575   |   |
| 3 | 3          | 3            | 3              | 7      | 1681       | 522   |   |
| 4 | 4          | 4            | 4              | 8      | 1375       | 250   |   |

|   | Aggressivity | Aggressiveness | Distance[km] | Fuel Consumed[L] | \ |
|---|--------------|----------------|--------------|------------------|---|
| 0 | 475.645508   | 5748.696207    | 9.225222     | 0.895646         |   |
| 1 | 269.540959   | 6372.966896    | 2.223611     | 0.441901         |   |
| 2 | 266.112770   | 5665.167276    | 6.323556     | 1.008958         |   |
| 3 | 552.585767   | 5891.611649    | 2.743222     | 0.344075         |   |

|   |                |             |              |               |
|---|----------------|-------------|--------------|---------------|
| 4 | 529.245807     | 5535.915731 | 4.683417     | 0.937469      |
|   | Fuel Rate[gpm] | Weight      | Displacement | SpeedAverage  |
| 0 | 0.097087       | 2500.0      | 1.8          | 46.423729     |
| 1 | 0.198731       | 3500.0      | 2.5          | 40.748170     |
| 2 | 0.159556       | 4000.0      | 2.4          | 32.929095     |
| 3 | 0.125427       | 3000.0      | 2.4          | 33.915000     |
| 4 | 0.200168       | 5500.0      | 5.3          | 33.196768     |
|   |                |             |              | SpeedVariance |
|   |                |             |              | 147.494133    |
|   |                |             |              | 447.766401    |
|   |                |             |              | 247.070941    |
|   |                |             |              | 447.276871    |
|   |                |             |              | 319.067811    |

```
[15]: data['SpeedAverage'].isna().sum()
```

```
[15]: 0
```

```
[16]: data['SpeedVariance'].isna().sum()
```

```
[16]: 0
```

## 1.1 Lets build a linear regression

```
[17]: # lets build a super simple stupid regression
from sklearn.linear_model import LinearRegression
```

```
[246]: # Before we can fit we have to do some preprocessing
features = ['Aggressiveness', 'Distance[km]', 'Weight', 'Displacement',
            ↪ 'SpeedAverage', 'SpeedVariance']
target = ['Fuel Rate[gpm]']
```

```
[127]: # WEIGHT HAS SOME NAN VALUES THAT ARE ENCODED AS FLOATS -- WE NEED TO USE NP.
        ↪ ISNAN INSTEAD OF PANDAS
#
mask = [~(np.isnan(data.iloc[i,:]).sum())>0) for i in range(len(data))]
data = data.loc[mask]
```

```
[212]: X = np.array(data.loc[:,features])
```

```
[213]: y = np.array(data.loc[:, 'Fuel Rate[gpm]']).reshape(-1,1)
```

```
[214]: X.shape
```

```
[214]: (11529, 6)
```

```
[215]: y.shape
```

```
[215]: (11529, 1)
```

```
[216]: model = LinearRegression().fit(X,y)
```

```
[217]: for feature, coef in zip(features, model.coef_[0]):
        print(f"{feature}: {coef}")
```

```
Aggressiveness: 1.188820516820224e-05
Distance[km]: -3.8248981102592604e-05
Weight: 6.059327624259096e-06
Displacement: 0.0299524660281594
SpeedAverage: 1.707837633178612e-05
SpeedVariance: 4.868811892222443e-08
```

```
[218]: # lets see which value has the highest average impact on our regression:

print('Impact = MEAN * Coef')
for i, vals in enumerate(zip(features, model.coef_[0])):
    feature, coef = vals
    print(f"{feature}: {X[:,i].mean()*coef}")

# we can actually get the relative percentages here:
print('===== PERCENTAGES =====')

impact = np.array([np.abs(coef * X[:,i].mean()) for i, coef in enumerate(model.
    ↪coef_[0])])
for feature, i in zip(features, (impact*100)/impact.sum()):
    print(f"{feature}: {i}")
```

```
Impact = MEAN * Coef
Aggressiveness: 0.06184528612055701
Distance[km]: -0.00020453144069008174
Weight: 0.020393534577180465
Displacement: 0.07506926104927253
SpeedAverage: 0.0006729971659087428
SpeedVariance: 2.4989540075156042e-05
===== PERCENTAGES =====
Aggressiveness: 39.090482029722686
Distance[km]: 0.12927796293517937
Weight: 12.890118987529737
Displacement: 47.44894532965449
SpeedAverage: 0.4253805790263045
SpeedVariance: 0.015795111131585857
```


```
[219]: # now lets compute r^2:
preds = model.predict(X)
r_2 = 1 - (((y - preds)**2).sum())/((y - y.mean())**2).sum()
```



```
[220]: r_2
```

```
[220]: 0.6123278886732875
```


```
[221]: # now lets draw our plots:
```

```
[224]: # single value predictions
for i, feature in enumerate(features):
    plt.scatter(X[:,i], y)
    plt.scatter(X[:,i], model.coef_[0][i]*X[:,i] + model.intercept_[0],
        ↪color='red')
    plt.title(f"Fuel Rate vs {feature} with predicted values")
    plt.show()
```







Fuel Rate vs Displacement with predicted values





Fuel Rate vs SpeedAverage with predicted values


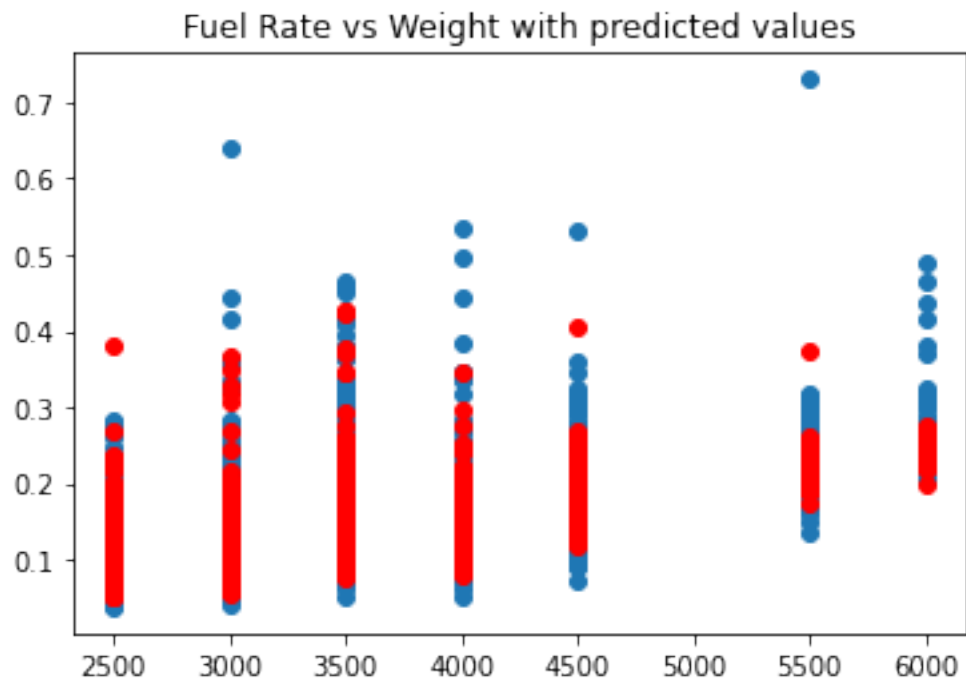






```
[225]: #plotting full model values
for i, feature in enumerate(features):
    plt.scatter(X[:,i], y)
    plt.scatter(X[:,i], preds, color='red')
    plt.title(f"Fuel Rate vs {feature} with predicted values")
    plt.show()
```










```
[226]: # lets plot our residuals:  
residuals = (preds - y)
```

```
plt.scatter(preds, residuals)
plt.title('Residuals vs fitted')
plt.show()
```



## 1.2 Adding Some Interaction Variables

```
[227]: # thats not a particularly compelling plot -- lets add some interaction
        ↪ variables to our set:
        for i, f in enumerate(features):
            print(f"{i} : {f}")
```

```
0 : Aggressiveness
1 : Distance[km]
2 : Weight
3 : Displacement
4 : SpeedAverage
5 : SpeedVariance
```

Lets interact the following variables:

Aggressiveness and SpeedAverage

Aggressiveness and Distance[km]

Displacement and  $1/\text{Weight}$  => power to weight

```
[228]: X.shape
```

```
[228]: (11529, 6)
```

```
[229]: X = np.append(X, (X[:,0] * X[:,4]).reshape(-1,1), 1) #aggr and avespeed  
X = np.append(X, (X[:,0] * X[:,3]).reshape(-1,1), 1) # aggr and distance  
X = np.append(X, (X[:,3] / X[:,2]).reshape(-1,1), 1) # displ and 1/weight
```

```
[230]: X.shape
```

```
[230]: (11529, 9)
```

```
[231]: #lets update our feature list  
interaction_features = features + ['Aggr*AveSpeed', 'Aggr*Distance', 'Displ/  
↪Weight']
```

```
[232]: # lets fit our model again:  
model = LinearRegression().fit(X,y)
```

```
[233]: for feature, coef in zip(interaction_features, model.coef_[0]):  
        print(f"{feature}: {coef}")
```

```
Aggressiveness: 4.109603391407518e-06  
Distance[km]: 1.7205819257853772e-05  
Weight: -9.172978419857352e-06  
Displacement: 0.03477552361547563  
SpeedAverage: -5.035984700663981e-05  
SpeedVariance: 1.9055986521698432e-07  
Aggr*AveSpeed: 1.2604850684488406e-08  
Aggr*Distance: 2.9460774284102315e-06  
Displ/Weight: -70.59826830183906
```

```
[234]: # lets see which value has the highest average impact on our regression:  
  
print('Impact = MEAN * Coef')  
for i, vals in enumerate(zip(interaction_features, model.coef_[0])):  
    feature, coef = vals  
    print(f"{feature}: {X[:,i].mean()*coef}")  
  
# we can actually get the relative percentages here:  
print('===== PERCENTAGES =====')  
  
impact = np.array([np.abs(coef * X[:,i].mean()) for i, coef in enumerate(model.  
↪coef_[0])])  
for feature, i in zip(interaction_features, (impact*100)/impact.sum()):  
    print(f"{feature}: {i}")
```

```

Impact = MEAN * Coef
Aggressiveness: 0.021379139574695276
Distance[km]: 9.200587570222754e-05
Weight: -0.030872972082272087
Displacement: 0.08715719293232776
SpeedAverage: -0.001984499793928708
SpeedVariance: 9.780627171411371e-05
Aggr*AveSpeed: 0.00258460756699648
Aggr*Distance: 0.038407058144407265
Displ/Weight: -0.052105052045205595
===== PERCENTAGES =====
Aggressiveness: 9.10989821095411
Distance[km]: 0.03920476591345401
Weight: 13.15532985583848
Displacement: 37.138686203526134
SpeedAverage: 0.8456182747292637
SpeedVariance: 0.04167638162403438
Aggr*AveSpeed: 1.1013311255270803
Aggr*Distance: 16.365690913579872
Displ/Weight: 22.202564268307555

```

```

[235]: # now lets compute r^2:
# WOW WE have an almost 0 impact on r^2 -- surprising
preds = model.predict(X)
r_2 = 1 - (((y - preds)**2).sum())/((y - y.mean())**2).sum()
print(r_2)

```


```
0.6200223135986707
```

```


[236]: # single value predictions
for i, feature in enumerate(interaction_features):
    plt.scatter(X[:,i], y)
    plt.scatter(X[:,i], model.coef_[0][i]*X[:,i] + model.intercept_[0],
        color='red')
    plt.title(f"Fuel Rate vs {feature} with predicted values")
    plt.show()



```



Fuel Rate vs Aggressiveness with predicted values





Fuel Rate vs Distance[km] with predicted values

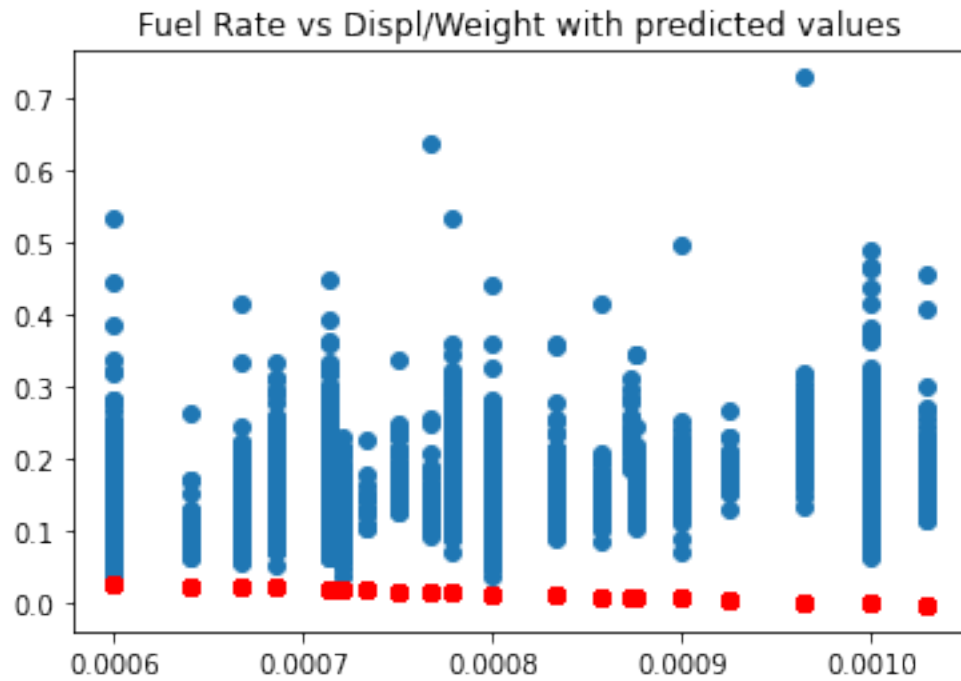




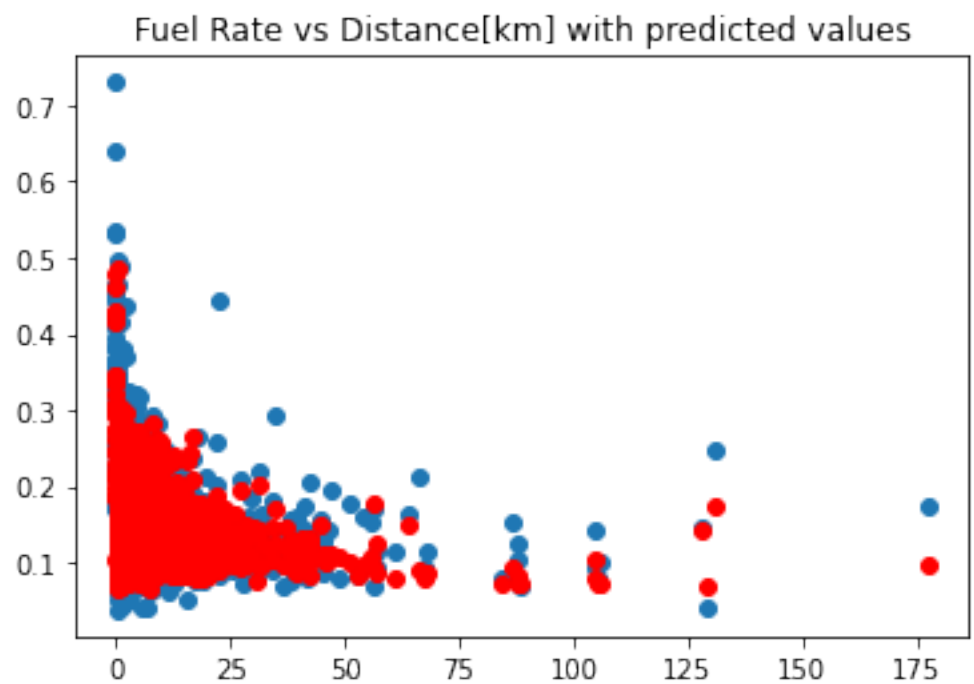








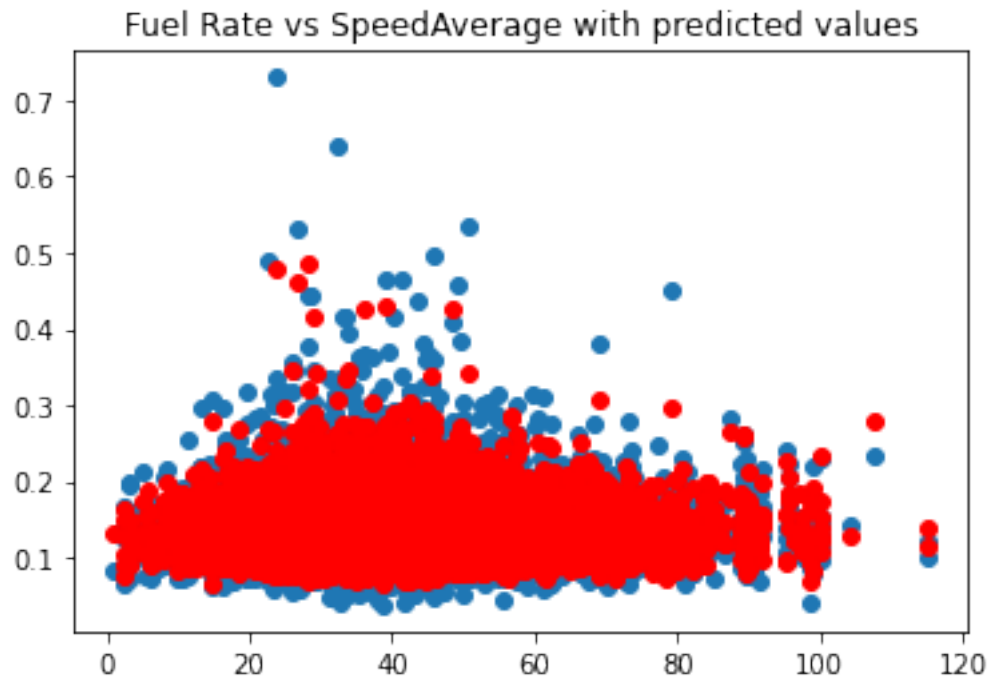





```
[238]: #plotting full model values
for i, feature in enumerate(interaction_features):
    plt.scatter(X[:,i], y)
    plt.scatter(X[:,i], preds, color='red')
    plt.title(f"Fuel Rate vs {feature} with predicted values")
    plt.show()
```









Fuel Rate vs Aggr\*AveSpeed with predicted values




Fuel Rate vs Aggr\*Distance with predicted values





```
[239]: # residual vs fitted plots!
residuals = (preds - y)
plt.scatter(preds, residuals)
plt.title('Residuals vs fitted')
plt.show() # we have the same issue with this plot where our residuals vs_
↳fitted are still clumpy
```



### 1.3 Lets run a larger interaction variable test:

```
[255]: # reset the data
X = np.array(data.loc[:, features])
y = np.array(data.loc[:, target]).reshape(-1,1)
```

```
[256]: X.shape
```

```
[256]: (11529, 6)
```

```
[257]: y.shape
```

```
[257]: (11529, 1)
```

```
[258]: # lets make all the interaction variables:
all_interaction = features.copy()
for i in range(len(features)):
    for j in range(i+1, len(features)):
        X = np.append(X, (X[:, i] * X[:, j]).reshape(-1,1), 1)
        all_interaction.append(f"{features[i]}*{features[j]}")
```

```
[259]: all_interaction
```



```
[259]: ['Aggressiveness',
        'Distance[km]',
        'Weight',
        'Displacement',
        'SpeedAverage',
        'SpeedVariance',
        'Aggressiveness*Distance[km]',
        'Aggressiveness*Weight',
        'Aggressiveness*Displacement',
        'Aggressiveness*SpeedAverage',
        'Aggressiveness*SpeedVariance',
        'Distance[km]*Weight',
        'Distance[km]*Displacement',
        'Distance[km]*SpeedAverage',
        'Distance[km]*SpeedVariance',
        'Weight*Displacement',
        'Weight*SpeedAverage',
        'Weight*SpeedVariance',
        'Displacement*SpeedAverage',
        'Displacement*SpeedVariance',
        'SpeedAverage*SpeedVariance']
```

```
[260]: X.shape
```

```
[260]: (11529, 21)
```

```
[261]: model = LinearRegression().fit(X, y)
```

```
[264]: for feature, coef in zip(all_interaction, model.coef_[0]):
        print(f"{feature}: {coef}")
```

```
Aggressiveness: 2.5022926173706823e-06
Distance[km]: 0.0003009805061247783
Weight: -1.1893202490620523e-05
Displacement: 0.010147303150106018
SpeedAverage: 6.144613174886634e-05
SpeedVariance: 1.974361128572301e-07
Aggressiveness*Distance[km]: -3.8251066330590243e-07
Aggressiveness*Weight: 1.7675447132677472e-09
Aggressiveness*Displacement: 1.6223401111833852e-06
Aggressiveness*SpeedAverage: 1.7993876257149626e-08
Aggressiveness*SpeedVariance: -6.893053204528441e-10
Distance[km]*Weight: 3.375073559130909e-07
Distance[km]*Displacement: -0.00022565392091916604
Distance[km]*SpeedAverage: 5.175361508629751e-06
Distance[km]*SpeedVariance: -1.7215678674891726e-09
Weight*Displacement: 3.2538600924069544e-06
```

```

Weight*SpeedAverage: -1.4078378054548045e-08
Weight*SpeedVariance: 4.779569972704856e-10
Displacement*SpeedAverage: -3.919495876954216e-05
Displacement*SpeedVariance: 1.8208014369773895e-06
SpeedAverage*SpeedVariance: -4.7788005972552274e-08

```

```

[265]: # lets see which value has the highest average impact on our regression:

print('Impact = MEAN * Coef')
for i, vals in enumerate(zip(all_interaction, model.coef_[0])):
    feature, coef = vals
    print(f"{feature}: {X[:,i].mean()*coef}")

# we can actually get the relative percentages here:
print('===== PERCENTAGES =====')

impact = np.array([np.abs(coef * X[:,i].mean()) for i, coef in enumerate(model.
    ↪coef_[0])])
for feature, i in zip(all_interaction, (impact*100)/impact.sum()):
    print(f"{feature}: {i}")

```

```

Impact = MEAN * Coef
Aggressiveness: 0.01301752457070437
Distance[km]: 0.0016094540236826912
Weight: -0.04002827562167602
Displacement: 0.02543198107311982
SpeedAverage: 0.002421370259072949
SpeedVariance: 0.00010133555708755586
Aggressiveness*Distance[km]: -0.009542428587891507
Aggressiveness*Weight: 0.031002679734935045
Aggressiveness*Displacement: 0.021149923073762494
Aggressiveness*SpeedAverage: 0.003689619964404585
Aggressiveness*SpeedVariance: -0.0018368910594713528
Distance[km]*Weight: 0.006168451356207922
Distance[km]*Displacement: -0.003084888067653157
Distance[km]*SpeedAverage: 0.0010893730424327756
Distance[km]*SpeedVariance: -4.766858874831378e-06
Weight*Displacement: 0.02881071700321241
Weight*SpeedAverage: -0.001866221008820074
Weight*SpeedVariance: 0.0008241984075872981
Displacement*SpeedAverage: -0.0038689221787572207
Displacement*SpeedVariance: 0.0023391801100916274
SpeedAverage*SpeedVariance: -0.0010369376461777547
===== PERCENTAGES =====
Aggressiveness: 6.543931361660887
Distance[km]: 0.8090752280530229
Weight: 20.122281065893784
Displacement: 12.784699397310199

```

```

SpeedAverage: 1.2172268767750087
SpeedVariance: 0.05094155393944858
Aggressiveness*Distance[km]: 4.796994802164125
Aggressiveness*Weight: 15.585098926530561
Aggressiveness*Displacement: 10.632101683186658
Aggressiveness*SpeedAverage: 1.8547781236391274
Aggressiveness*SpeedVariance: 0.9234082061255259
Distance[km]*Weight: 3.100890807889174
Distance[km]*Displacement: 1.550778388278212
Distance[km]*SpeedAverage: 0.5476296494163669
Distance[km]*SpeedVariance: 0.0023963079246127894
Weight*Displacement: 14.483195597229951
Weight*SpeedAverage: 0.9381524206907869
Weight*SpeedVariance: 0.414325917216198
Displacement*SpeedAverage: 1.9449136465140398
Displacement*SpeedVariance: 1.1759097514938426
SpeedAverage*SpeedVariance: 0.5212702880684681

```

```

[266]: # now lets compute r^2:
# WOW WE have an almost 0 impact on r^2 -- surprising
preds = model.predict(X)
r_2 = 1 - (((y - preds)**2).sum())/((y - y.mean())**2).sum()
print(r_2)

```

```
0.634915533525551
```

## 1.4 Optimizing our Linear Regression with Smart Variable Choices

```
[ ]:
```